

Feature Engineering

Learning Objectives

- Introduction to Feature engineering
- Hands on exercise on Feature engineering
- Part 2 of feature engineering and model tuning
- How to tune the models or improve performance
- Concept of upsampling and downsampling
- Hands on exercise showing tuning of a model

Introduction to Feature engineering

- **Feature engineering** is the process of using domain knowledge of the data to create features that make machine learning algorithms work.
- It is fundamental to the application of machine learning, and is both difficult and expensive.
- One reason why feature engineering is so important is that defining and/or learning higher level domain specific feature is actually one way to deep learning.

Hands on exercise on Feature engineering

Some important functions

1. Importing packages

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.model_selection import train_test_split
```

```
From sklearn import metrics
```

```
from sklearn.preprocessing import Imputer
```

```
from sklearn.tree import DecisionTreeClassifier
```

2. Split dataset into inputs and outputs

```
values = pima_df.values
```

```
X = values[:,0:8]
```

```
Y = values[:,8]
```

```
test_size = 0.30 # taking 70:30 training and test set
```

3. Fill missing values with mean column values

```
imputer = Imputer(missing_values = 'NaN', strategy='median', axis=0)
```

```
transformed_x= imputer.fit_transform(X)
```

Part 2 of feature engineering and model tuning

- Feature engineering is a set of those activities which performs a discipline we follow, to ensure that you finally feed the right data to your algorithms, in which the target variable is very strongly influenced by the independent variables.
- The independent variables do not have any significant outliers, all the missing values have been taken care of.
- As a part of data preparation, we have 2 techniques such as Upsampling and Downsampling

Hands on exercise feature engineering part 2

Analysing car mpg data set using clustering and PCA

Some important functions

1. Importing libraries

```
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LinearRegression
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
```

2. Drop unique columns

```
Numeric_cols = mpg_df.drop('car_name', axis=1)
```

3. Pair plot analysis to visually check number of likely clusters

```
mpg_df_attr = mg_df.iloc[:,0:9]
```

```
sns.pairplot(mpg_df_attr, diag_kind = 'kde') #to plot density curve instead of histogram
```

4. Splitting the data

```
array = mpg_df_attr.values
```

```
X= array[:, 1:5]
```

```
Y= array[:, 0]
```

```
X_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.30, random_state=1)
```


Iteration 2

5. Drop acc column and follow the same steps for splitting

```
mpg_df_attr_z.pop('acc')
```

```
array = mpg_df_attr_z.values
```

And same splitting steps

6. Create a regularized RIDGE model and note the coefficients

```
ridge= Ridge(alpha = 0.3)
```

```
ridge.fit(x_train, y_train)
```

```
print("Ridge model:", (ridge.coef_))
```

```
Ridge model: [[ 2.47057456  2.4449775  ...]]
```

7. Create a regularized LASSO model and note the coefficients

```
lasso= Lasso(alpha = 0.2)
lasso.fit(x_train, y_train)
print("Lasso model:" , (lasso.coef_))
Lasso model: [[0.  0. -0.3456 -4.09876 ---]]
```

8. Generate polynomial models reflecting the non-linear interaction between some dimensions


```
from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(degree = 2, interaction_only = True)
X_poly = poly.fit_transform(X_scaled)
X_train, x_test, y_train, y_test = train_test_split (x_poly, y, test_size=0.30, random_state = 1)
X_train.shape
```

How to tune the models or improve performance

- To increase performance of any model, increase model capacity.
- The tuning process is more empirical than theoretical. We add layers and nodes gradually with the intention to overfit the model since we can tone it down with regularizations.
- We repeat the iterations until the accuracy improvement is diminishing and no longer justify the drop in the training and computation performance

Concept of upsampling and downsampling

- We can handle the imbalanced dataset cases to minimize the Type II errors by balancing the class representations.
- To balance the classes we can –
 - Decrease the frequency of the majority class
 - Increase the frequency of the minority class



Imblearn techniques

- Python imbalanced-learn module – provides more sophisticated resampling techniques.
- In over-sampling, instead of creating exact copies of the minority class records, we can introduce small variations into those copies creating more diverse synthetic samples.
- SMOTE consists of synthesizing elements for the minority class, based on those that already exist, works randomly picking a point from the majority class and computing the K-nearest neighbours.

Hands on exercise showing tuning of a model

1. Importing libraries

```
import pandas as pd  
import matplotlib.pyplot as plt  
import matplotlib.style  
plt.style.use('classic')  
import seaborn as sns
```

2. Copy all the predictor variables into X dataframe.

```
X= mpg_df.drop('mpg', axis=1)  
X= x.drop({'origin_america', 'origin_asia', 'origin_europe'}, axis=1)
```

3. Prepare the model after splitting

```
regression_model = LinearRegression()  
regression_model.fit(x_train, y_train)
```

4. Explore the coefficients for each of the independent attributes

```
for idx, col_name in enumerate(x_train.columns):  
print("The coefficient for {} is {}".format(col_name, regression_model.coef_[col_names]))
```

5. Down sampling the larger class i.e. non diabetics

```
non_diab_indices = pima_df[pima_df['class'] == 0].index  
no_diab = len(pima_df['class'] == 0)  
print(no_diab)
```

```
pima_df_down_sample = pima_df.loc [down_sample_indices] #extract all those recommendations  
pima_df_down_sample.shape  
pima_df_down_sample.groupby(["class"]).count # Look at the class distribution
```

6. Fit the model on 30% through Logistic Regression

```
model = LogisticRegression()  
model.fit(x_train, y_train)  
y_predict = model.predict(x_test)  
model_score = model.score(x_test, y_test)  
print(model_score)  
print(metrics.confusion_matrix(y_test, y_predict))
```




Questions?

